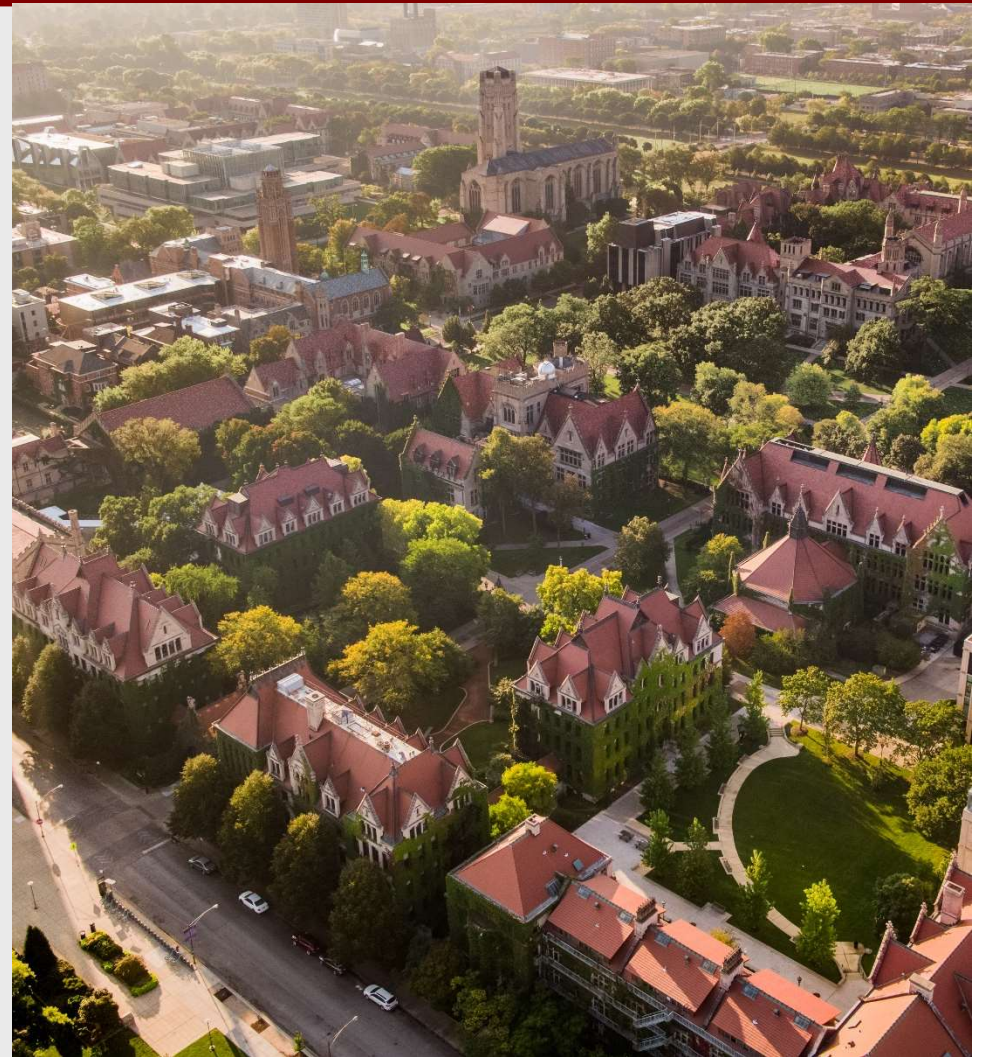


---

# University of Chicago Professional Education

MSCA 37010 Programming for Analytics  
Week 3 Lecture Notes

Autumn 2020



## ❑ Introduction

Instructor: Mei Najim

Email: [mnajim@uchicago.edu](mailto:mnajim@uchicago.edu)

Tel: 847-800-9979 (C)

Class Meeting Time: 6:00 - 9:00pm, Mondays (01 Section)

1:30 - 4:30pm Saturdays (02 Section)

Tentative Office Hour: Mondays (9:00pm – 10:00pm or until last students)

Wednesdays (5:30pm – 6:30pm or until last students)

Saturdays (4:30pm – 5:30pm or until last students)

Notes: 1) First ten-minute quiz; Breakout groups in zoom; Two 10-minute breaks

2) Set up a weekly discussion group on canvas, allow 24 hours to respond

3) Email questions with the **Section Number (01/02)** in the subject line

4) If it is urgent, feel free to text me directly (847-800-9979)

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Subsetting

- Subsetting vectors, matrices, arrays, or data frame can also be referred to as indexing
  - 1) Subsetting Vectors
  - 2) Subsetting Data Frames
  - 3) Subsetting Operators in R

## ❑ Data Preparation – Subsetting Vectors

- Subsetting Vectors:

- ✓ Subsetting the elements of a vector can be achieved by inserting an index vector in **square brackets** to the name of the vector

- ✓ A logical vector:

- The index vector should be the same length as the vector from which the elements are to be selected
    - Values corresponding to TRUE in index vector are being selected / to FALSE are being excluded
    - Values corresponding to NA returns NA

## ❑ Data Preparation – Subsetting Vectors

- Subsetting Vectors:

Examples:

```
> a <- c(1, 3, 5, NA, 7)
```

```
> is.na(a)    [1] FALSE FALSE FALSE TRUE FALSE
```

```
> !is.na(a)   [1] TRUE TRUE TRUE FALSE TRUE
```

```
> a[!is.na(a)] [1] 1 3 5 7
```

```
> a > 3       [1] FALSE FALSE TRUE NA TRUE
```

```
> a[a > 3]    [1] 5 NA 7
```

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames

- Data frame shares the common properties of matrices and lists. If you subset with two vectors at both positions, it behaves like subsetting a matrix

```
> d <- data.frame(L = c("A", "B", "C"), M = 1:3, N = c(T, F, NA))
> d[c(2,3), ]
  L M    N
2 B 2 FALSE
3 C 3    NA
> d[,c(1,3)]
  L    N
1 A  TRUE
2 B FALSE
3 C    NA
> d[d$L == "A", ]
  L M    N
1 A 1  TRUE
> d[, c("M", "N")]
  M    N
1 1  TRUE
2 2 FALSE
3 3    NA
> d[d$L == "A", c("M", "N")]
  M    N
1 1  TRUE
```



## ❑ Data Preparation – Subsetting Operators

- Subsetting Operators: **Three operators, “[ ]”, “[[ ]]”, and “\$”, can be used to subset objects.** Deciding which operator to use depends upon the object type. The main behavior difference across these three objects are:
  - The type of object returned by using the “[“ operator is the same data type as the object “[“ applies to.  
Example: using “[“ to subset a list returns a list
  - You can use “[“ to extract any numbers of elements of an object, while you can only use “[[“ and “\$” to extract one element
  - **“\$” does not evaluate its argument, while “[[“ and “[“ do.** Thus, you can include an expression inside “[[“ or “[“
  - **“\$” uses partial matching to extract elements, while “[[“ and “[“ do not**

## ❑ Data Preparation – Subsetting Operators

- Examples

```
> alist <- list(name1 = c("john", "ken"), station = "AM640", time = "M-F: 3:00pm")
```

```
> alist[c(1,2)]
```

```
$`name1` [1] "john" "ken"
```

```
$station [1] "AM640"
```

- The code above extracts the first two elements from a list. The resulting object is also a list. You can not write `alist[[c(1,2)]]` since you can **only extract one element by using the `[[` operator**

```
> alist["name1"]    $`name1` [1] "john" "ken"
```

```
> alist[["name1"]]  [1] "john" "ken"
```

```
> alist$name1       [1] "john" "ken"
```

## ❑ Data Preparation – Subsetting Data Frames

- Simplifying : using the index vector with “[” operator at the second index position or using the “[[” operator when selecting a column
- Preserving : using the index vector the “[” operator at the second index position and setting the drop option to FALSE or using the “[[” operator when selecting a column

```
> dat <- data.frame(V1=1:3, V2=c("a", "b", "c"), V3=c(T, T, F)); dat
  V1 V2  V3
1  1  a TRUE
2  2  b TRUE
3  3  c FALSE
> dat[, 2]
[1] a b c
Levels: a b c
> dat[[2]]
[1] a b c
Levels: a b c
> dat[, 2, drop = F]
  V2
1  a
2  b
3  c
> dat[2]
  V2
2  b
3  c
```

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames by Using Index Vectors

We will use the data set, painters, from the MASS library to illustrate some examples

```
> library(MASS)
```

```
> head(painters)
```

	Composition	Drawing	Colour	Expression	School
Da Udine	10	8	16	3	A
Da Vinci	15	16	4	14	A
Del Piombo	8	13	16	7	A
Del Sarto	12	16	9	8	A
Fr. Penni	0	15	8	0	A
Guilio Romano	15	16	4	14	A

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames by Using Index Vectors

Selecting the observations from a data frame is similar to selecting rows from a matrix by placing an index vector on the left side of the comma (,) inside the []

Example: To select the observations of painters with Colour greater than or equals 17:

```
> painters[painters$Colour>=17,]
```

	Composition	Drawing	Colour	Expression	School
Bassano	6	8	17	0	D
Giorgione	8	9	18	4	D
Pordenone	8	14	17	5	D
Titian	12	15	18	6	D
Rembrandt	15	6	17	12	G

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames by Using Index Vectors

Example: To select those from school A and D, you may want to use `School == c('A', 'D')`. This tests

`> painters[painters$School %in% c('A', 'D'),]`

	Composition	Drawing	Colour	Expression	School
Da Udine	10	8	16	3	A
Da Vinci	15	16	4	14	A
Del Piombo	8	13	16	7	A
Del Sarto	12	16	9	8	A
Fr. Penni	0	15	8	0	D

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames by Using Index Vectors

Selecting variables from a data frame is also similar to selecting columns from a matrix by placing an index vector on the right side of the comma (,) inside the []

Example: To create a data set that contain the Colour (the third column) and School variables (the fifth column), you can write either one of the following statements

```
> d1 <- painters[, c('School', 'Colour')]
```

```
> d2 <- painters[, c(5,3)]
```

```
> head(d2)
```

	School	Colour
Da Udine	A	16
Da Vinci	A	4
Del Piombo	A	16

## ❑ Data Preparation – Subsetting Data Frames

- Subsetting Data Frames by Using Index Vectors

You can also select observations and variables at the same time by including two index vectors inside

```
> d5<-painters[painters$School == "A", c('School', 'Colour')]; d5
```

	School	Colour
Da Udine	A	16
Da Vinci	A	4
Del Piombo	A	16
Del Sarto	A	9
Fr. Penni	A	8
Guilio Romano	A	4
Michelangelo	A	4



## ❑ Data Preparation – Subsetting Data Frames Example

- Subsetting Data Frames by Using *subset()* function

Examples: Use AutoCollision data

- Subsetting ClaimSeverity>=300 **and** ClaimCount>=10

```
> subset(AutoColl, ClaimSeverity >=200 & ClaimCount>=10);
```

- Subsetting ClaimSeverity>=300 **or** ClaimCount>=10

```
> subset(AutoColl, ClaimSeverity >=200 | ClaimCount>=10);
```

- Subsetting ClaimSeverity >=300 without AgeGroup column

```
> AutoColl[AutoColl$ClaimSeverity>=300, c('ClaimSeverity','VehicleUse')]
```

```
> subset(AutoColl, ClaimSeverity >=300, select= - AgeGroup)
```

```
> subset.data.frame(AutoColl, ClaimSeverity >=300, select= - AgeGroup)
```

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Merging

- Merging: Merge two datasets by an ID variable, where ID is the same for both datasets

```
> # Merge two datasets by an ID variable, where ID is the same for both datasets
> data1 <- data.frame(ID=1:5, x=letters[1:5]);data1
  ID x
1  1 a
2  2 b
3  3 c
4  4 d
5  5 e
> data2 <- data.frame(ID=1:5, y=letters[6:10]);data2
  ID y
1  1 f
2  2 g
3  3 h
4  4 i
5  5 j
> data3<-merge(data1, data2);data3
  ID x y
1  1 a f
2  2 b g
3  3 c h
4  4 d i
5  5 e j
```

## ❑ Data Preparation – Merging

- Merging: Merge two datasets by an ID variable, where ID is not the same for both datasets

```
> data1 <- data.frame(ID=1:5, x=letters[1:5])
> data2 <- data.frame(ID=4:8, y=letters[6:10])
> merge(data1, data2)
  ID x y
1  4 d f
2  5 e g
> merge(data1, data2, all=TRUE)
  ID x   y
1  1 a <NA>
2  2 b <NA>
3  3 c <NA>
4  4 d   f
5  5 e   g
6  6 <NA>  h
7  7 <NA>  i
8  8 <NA>  j
> merge(data1, data2, all.x=TRUE) # Only keep the rows from the 1st argument data1
  ID x   y
1  1 a <NA>
2  2 b <NA>
3  3 c <NA>
4  4 d   f
5  5 e   g
> merge(data1, data2, all.y=TRUE) # Only keep the rows from the 2nd argument data2
  ID x y
1  4 d f
2  5 e g
3  6 <NA> h
4  7 <NA> i
5  8 <NA> j
```

## ❑ Data Preparation – Merging

- Merging: Merge two datasets by an ID variable, where both dataset have the same names

```
> data1 <- data.frame(ID=1:5, x=letters[1:5])
> data2 <- data.frame(ID=1:5, x=letters[6:10])
> merge(data1, data2, all=TRUE) # Add rows
  ID x
1  1 a
2  1 f
3  2 b
4  2 g
5  3 c
6  3 h
7  4 d
8  4 i
9  5 e
10 5 j
> merge(data1, data2, by="ID") # Add columns
  ID x.x x.y
1  1  a  f
2  2  b  g
3  3  c  h
4  4  d  i
5  5  e  j
> merge(data1, data2, by="ID", suffixes=c(1, 2))
  ID x1 x2
1  1  a  f
2  2  b  g
3  3  c  h
4  4  d  i
5  5  e  j
```

## ❑ Data Preparation – Merging

- Merging: Merge two datasets by an ID variable, where the ID variable has a different name

```
> data1 <- data.frame(ID1=1:5, x=letters[1:5]);data1
  ID1 x
1   1 a
2   2 b
3   3 c
4   4 d
5   5 e
> data2 <- data.frame(ID2=1:5, x=letters[6:10]);data2
  ID2 x
1   1 f
2   2 g
3   3 h
4   4 i
5   5 j
> merge(data1, data2, by.x="ID1", by.y="ID2")
  ID1 x.x x.y
1   1  a  f
2   2  b  g
3   3  c  h
4   4  d  i
5   5  e  j
```

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Reviewing Data Types: Numeric

- R Technically, numeric data in R can be either double or integer, but in practice numeric data is almost always double (type double refers to real numbers). See `?integer` and `?double`
- ***format()*** formats an object for pretty printing. `format()` is a generic function that is used with other types of objects. See `?format()` for additional arguments.

# trim - If FALSE right justified with common width

```
> format(c(1,10,100,1000), trim = FALSE) [1] " 1" " 10" " 100" "1000"
```

```
> format(c(1,10,100,1000), trim = TRUE) [1] "1" "10" "100" "1000"
```

# nsmall - Minimum number of digits to the right of the decimal point

```
format(13.7, nsmall = 3) [1] "13.700"
```

# scientific - Use scientific notation

```
> format(2^16, scientific = TRUE) [1] "6.5536e+04"
```



## ❑ Data Types – Numeric: Integer vs. Double

- The two most common numeric classes used in R are integer and double (for double precision floating point numbers). R automatically converts between these two classes when needed for mathematical purposes. As a result, it's feasible to use R and perform analyses for years without specifying these differences.
- Creating Integer and Double Vectors: By default, when you create a numeric vector using the `c()` function it will produce a vector of double precision numeric values. To create a vector of integers using `c()` you must specify explicitly by placing an `L` directly after each number.

# create a string of double-precision values

```
> dbl_var <- c(1, 2.5, 4.5) ; dbl_var      [1] 1.0 2.5 4.5
```

# placing an L after the values creates a string of integers

```
> int_var <- c(1L, 6L, 10L); int_var      [1] 1 6 10
```

## ❑ Data Types – Numeric: Integer vs. Double

- The two most common numeric classes used in R are integer and double (for double precision floating point) Checking for Numeric Type

To check whether a vector is made up of integer or double values:

# identifies the vector type (double, integer, logical, or character)

```
> typeof(dbl_var) [1] "double"
```

```
> typeof(int_var) [1] "integer"
```

- Converting Between Integer and Double Values

By default, if you read in data that has no decimal points or you create numeric values using the `x <- 1:10` method the numeric values will be coded as integer. If you want to change a double to an integer or vice versa you can specify one of the following:

# converts integers to double-precision values

```
> as.double(int_var) [1] 1 6 10 # identical to as.double()
```

```
> as.numeric(int_var) [1] 1 6 10 # converts doubles to integers
```

```
> as.integer(dbl_var) [1] 1 2 4
```

## ❑ Data Types – Logical

- Logical values are represented by the reserved words **TRUE** and **FALSE** in all caps or simply **T** and **F**

<code>!x</code>	NOT <code>x</code>
<code>x &amp; y</code>	<code>x</code> AND <code>y</code> elementwise, returns a vector
<code>x &amp;&amp; y</code>	<code>x</code> AND <code>y</code> , returns a single value
<code>x   y</code>	<code>x</code> OR <code>y</code> elementwise, returns a vector
<code>x    y</code>	<code>x</code> OR <code>y</code> , returns a single value
<code>xor(x,y)</code>	Exclusive OR of <code>x</code> and <code>y</code> , elementwise
<code>x %in% y</code>	<code>x</code> IN <code>y</code>
<code>x &lt; y</code>	<code>x</code> < <code>y</code>
<code>x &gt; y</code>	<code>x</code> > <code>y</code>
<code>x &lt;= y</code>	<code>x</code> ≤ <code>y</code>
<code>x &gt;= y</code>	<code>x</code> ≥ <code>y</code>
<code>x == y</code>	<code>x</code> = <code>y</code>
<code>x != y</code>	<code>x</code> ≠ <code>y</code>
<code>isTRUE(x)</code>	TRUE if <code>x</code> is TRUE
<code>all(...)</code>	TRUE if all arguments are TRUE
<code>any(...)</code>	TRUE if at least one argument is TRUE
<code>identical(x,y)</code>	Safe and reliable way to test two objects for being <i>exactly</i> equal
<code>all.equal(x,y)</code>	Test if two objects are <i>nearly</i> equal

## ❑ Data Types – Example Logical Operations

Example:

```
> x <- 1:10;x
[1] 1 2 3 4 5 6 7 8 9 10
> (x%%2==0) | (x > 5) # What elements of x are even or greater than 5?
[1] FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE
> x[(x%%2==0) | (x > 5)]
[1] 2 4 6 7 8 9 10
> y <- 5:15 # What elements of x are in y? > x %in% y
[1] FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
> x[x%in% y]
[1] 5 6 7 8 9 10
> any(x>5) # Are any elements of x greater than 5?
[1] TRUE
> all(x>5) # Are all the elements of x greater than 5?
[1] FALSE
```

## ❑ Data Types – Isn't That Equal?

- In general, logical operators may not produce a single value and may return an **NA** if an element is **NA** or **NaN**.
- If you must get a single **TRUE** or **FALSE**, such as with **if** expressions, you should NOT use **==** or **!=**. Unless you are absolutely sure that nothing unusual can happen, you should use the ***identical()*** function instead.
- ***identical()*** only returns a single logical value, **TRUE** or **FALSE**, never **NA**

```
> name <- "Nick"; > if(name=="Nick") TRUE else FALSE      [1] TRUE
```

# But what if name is never set to "Nick"?

```
> name <- NA;      > if(name=="Nick") TRUE else FALSE
```

Error in if (name == "Nick") TRUE else FALSE : missing value where TRUE/FALSE needed

```
> if(identical(name, "Nick")) TRUE else FALSE            [1] FALSE
```

## ❑ Data Types – Isn't That Equal?

- With *all.equal()* objects are treated as equal if the only difference is probably the result of inexact floating-point calculations. Returns *TRUE* if the mean relative difference is less than the specified tolerance.
- *all.equal()* either returns TRUE or a character string that describes the difference. Therefore, do not use *all.equal()* directly in if expressions, instead use with *isTRUE()* or *identical()*.

```
> (x <- sqrt(2))
```

```
[1] 1.414214
```

```
> x^2
```

```
[1] 2
```

```
> x^2==2
```

```
[1] FALSE
```

```
> all.equal(x^2, 2)
```

```
[1] TRUE
```

```
> all.equal(x^2, 1)
```

```
[1] "Mean relative difference: 0.5"
```

```
> isTRUE(all.equal(x^2, 1))
```

```
[1] FALSE
```

## ❑ Data Types – Character

- Character strings are defined by quotation marks, single ' ' or double " "

cat()	Concatenate objects and print to console (\n for newline)
paste()	Concatenate objects and return a string
print()	Print an object
substr()	Extract or replace substrings in a character vector
strtrim()	Trim character vectors to specified display widths
strsplit()	Split elements of a character vector according to a substring
grep()	Search for matches to a pattern within a character vector, returns a vector of the indices that matched
grepl()	Like grep(), but returns a logical vector
agrep()	Similar to grep(), but searches for approximate matches
regexpr()	Similar to grep(), but returns the position of the first instance of a pattern <i>within</i> a string
gsub()	Replace all occurrences of a pattern with a character vector
sub()	Like gsub(), but only replaces the first occurrence
tolower(), toupper()	Convert to all lower/upper case
noquote()	Print a character vector without quotations
nchar()	Number of characters
letters, LETTERS	Built-in vector of lower and upper case letters

## ❑ Data Types – Example Character Functions

- Character strings are defined by quotation marks, single ' ' or double " “

```
> animals <- c("bird", "horse", "fish"); home <- c("tree", "barn", "lake")
```

```
> length(animals)          # Number of strings
```

R Output: [1] 3

```
> nchar(animals)           # Number of characters in each string
```

R Output: [1] 4 5 4

```
> cat("Animals:", animals)  # Need \n to move cursor to a newline
```

R Output: Animals: bird horse fish

```
> cat(animals, home, "\n")  # Joins one vector after the other
```

R Output: bird horse fish tree barn lake

```
> paste(animals, collapse=" ") # Create one long string of animals
```

R Output: [1] "bird horse fish"

```
> substr(animals, 2, 4)      # Get characters 2-4 of each animal
```

R Output: [1] "ird" "ors" "ish"

```
> strtrim(animals, 3)        # Print the first three characters
```

R Output: [1] "bir" "hor" "fis"

```
> toupper(animals)           # Print animals in all uppercase
```

R Output: [1] "BIRD" "HORSE" "FISH"



## ❑ Data Types – Factors

- A factor is a categorical variable with a defined number of ordered/unordered levels. Use the function `factor` to create a factor variable.

```
> factor(rep(1:2, 4), labels=c("trt.1", "trt.2")) [1] trt.1 trt.2 trt.1 trt.2 trt.1 trt.2 trt.1 trt.2
```

Levels: trt.1 trt.2

```
> factor(rep(1:3, 3), labels=c("low", "med", "high"), ordered=TRUE) [1] low med high low med high low med high
```

Levels: low < med < high

<code>levels(x)</code>	Retrieve or set the levels of <code>x</code>
<code>nlevels(x)</code>	Return the number of levels of <code>x</code>
<code>relevel(x, ref)</code>	Levels of <code>x</code> are reordered so that the level specified by <code>ref</code> is first
<code>reorder()</code>	Reorders levels based on the values of a second variable
<code>gl()</code>	Generate factors by specifying the pattern of their levels
<code>cut(x, breaks)</code>	Divides the range of <code>x</code> into intervals (factors) determined by <code>breaks</code>

## ❑ Data Types – Dates and Times

- R has objects that are dates only and objects that are dates and times. We will just focus on dates. Look at [?DateTimeClasses](#) for information about how to handles dates and times.
- An R date object has the format: Year-Month-Day
- Operations with dates:
  - Days can be added or subtracted to a date
  - Dates can be subtracted
  - Dates can be compared using logical operators

<code>Sys.Date()</code>	Current date
<code>as.Date()</code>	Convert a character string to a date object
<code>format.Date()</code>	Change the format of a date object
<code>seq.Date()</code>	Generate sequence of dates
<code>cut.Date()</code>	Cut dates into intervals
<code>weekdays, months, quarters</code>	Extract parts of a date object
<code>julian</code>	Number of days since a given origin

`.Date` suffix is optional for calling `format.Date()`, `seq.Date()` and `cut.Date()`, but is necessary for viewing the appropriate documentation

## ❑ Data Types – Dates and Times

- Converting a string to a date object requires specifying a format string that defines the date format
- Any character in the format string other than the % symbol is interpreted literally.
- Common conversion specifications (see ?strptime for a complete list),
  - %a Abbreviated weekday name
  - %A Full weekday name
  - %d Day of the month
  - %B Full month name
  - %b Abbreviated month name
  - %m Numeric month (01-12)
  - %y Year without century (2 digits)
  - %Y Year with century (4 digits)

Example:

```
> dates.1 <- c("5jan2008", "19aug2008", "2feb2009", "29sep2009")
> as.Date(dates.1, format="%d%b%Y")      [1] "2008-01-05" "2008-08-19" "2009-02-02" "2009-09-29"
> dates.2 <- c("5-1-2008", "19-8-2008", "2-2-2009", "29-9-2009")
> as.Date(dates.2, format="%d-%m-%Y")    [1] "2008-01-05" "2008-08-19" "2009-02-02" "2009-09-29"
```

## ❑ Data Types – Sequence of Dates

- To create a sequence of dates, **seq.Date(from, to, by, length.out = NULL)**
  - **from, to** Start and ending date objects
  - **by** A character string, containing one of "day", "week", "month" or "year". Can optionally be preceded by a (positive or negative) integer and a space, or followed by a "s".
  - **length.out** Integer, desired length of the sequence

Example:

```
> seq.Date(as.Date("2011/1/1"), as.Date("2011/1/31"), by="week")
[1] "2011-01-01" "2011-01-08" "2011-01-15" "2011-01-22" "2011-01-29"

> seq.Date(as.Date("2011/1/1"), as.Date("2011/1/31"), by="3 days")
[1] "2011-01-01" "2011-01-04" "2011-01-07" "2011-01-10" "2011-01-13" "2011-01-16" "2011-01-19"
[8] "2011-01-22" "2011-01-25" "2011-01-28" "2011-01-31"

> seq.Date(as.Date("2011/1/1"), by="week", length.out=10)
[1] "2011-01-01" "2011-01-08" "2011-01-15" "2011-01-22" "2011-01-29" "2011-02-05" "2011-02-12"
[8] "2011-02-19" "2011-02-26" "2011-03-05"
```

## ❑ Data Types – Cutting Dates

- To divide a sequence of dates into levels, `cut.Date(x, breaks, start.on.monday = TRUE)`

```
> jan <- seq.Date(as.Date("2011/1/1"), as.Date("2011/1/31"), by="days")
```

```
> cut(jan, breaks="weeks", start.on.monday=TRUE)
```

```
[1] 2010-12-27 2010-12-27 2011-01-03 2011-01-03 2011-01-03 2011-01-03 2011-01-03 2011-01-03 2011-01-03 2011-01-10
```

```
[11] 2011-01-10 2011-01-10 2011-01-10 2011-01-10 2011-01-10 2011-01-10 2011-01-17 2011-01-17 2011-01-17 2011-01-17
```

```
[21] 2011-01-17 2011-01-17 2011-01-17 2011-01-24 2011-01-24 2011-01-24 2011-01-24 2011-01-24 2011-01-24 2011-01-24
```

```
[31] 2011-01-31
```

```
Levels: 2010-12-27 2011-01-03 2011-01-10 2011-01-17 2011-01-24 2011-01-31
```

January 2011

Sun	Mon	Tue	Wed	Thr	Fri	Sat
26	27	28	29	30	31	1
2	3	4	5	6	7	8
9	10	11	12	13	14	15
16	17	18	19	20	21	22
23	24	25	26	27	28	29
30	31	1	2	3	4	5

## ❑ Data Types – Operations with Dates

- Operations with dates:
  - Days can be added or subtracted to a date
  - Dates can be subtracted
  - Dates can be compared using logical operators

Example:

```
> jan1 <- as.Date("2011/1/1")
> (jan8 <- jan1 + 7)           # Add 7 days to 2011/1/1           [1] "2011-01-08"
> jan1 - 14                   # Subtract 2 weeks from 2011/1/8   [1] "2010-12-18"
> jan8 - jan1                 # Number of days between 2011/1/1 and 2011/1/8 Time difference of 7 days
> jan8 > jan1                 # Compare dates                     [1] TRUE
# Use format to extract parts of a date object or change the appearance
> format.Date(jan8, "%Y")      [1] "2011"
> format.Date(jan8, "%b-%d")   [1] "Jan-08"
```

## ❑ Data Types – Testing and Coercing Objects

- All objects in R have a type. We can test the type of an object using a *is.type()* function.
- We can also attempt to **coerce** objects of one type to another using a *as.type()* function.
- Automatic conversions:
  - Logical values are converted to numbers by setting **FALSE** as **0** and **TRUE** as **1**
  - Logical, numeric, factor and date types are converted to characters by converting each element/level individually
- Some general rules for coercion:
  - Numeric values are coerced to logical by treating all **non-zero values** as **TRUE**
  - Numeric characters can be coerced to numbers, but non-numeric characters cannot
  - Factor levels can be coerced to numeric and numbers can be coerced to factors with a level for each unique number
  - Vectors, matrices and arrays are coerced to lists by making each element a vector of length 1
  - Vectors, matrices, arrays can also be coerced from one form to another

## ❑ Data Types – Testing and Coercing Functions

Type	Testing	Coercing
Array	<code>is.array()</code>	<code>as.array()</code>
Character	<code>is.character()</code>	<code>as.character()</code>
Dataframe	<code>is.data.frame()</code>	<code>as.data.frame()</code>
Factor	<code>is.factor()</code>	<code>as.factor()</code>
List	<code>is.list()</code>	<code>as.list()</code>
Logical	<code>is.logical()</code>	<code>as.logical()</code>
Matrix	<code>is.matrix()</code>	<code>as.matrix()</code>
Numeric	<code>is.numeric()</code>	<code>as.numeric()</code>
Vector	<code>is.vector()</code>	<code>as.vector()</code>



## ❑ Data Types – Testing and Coercing Objects

Example:

```
> x <- 1:4; > x>3
[1] FALSE FALSE FALSE TRUE
> sum(x>3)      # Automatic conversion to numeric vector; note TRUE=1, FALSE=0
[1] 1
> is.vector(x)
[1] TRUE
> is.numeric(x)
[1] TRUE
> as.list(x)
[[1]]
[1] 1
[[2]]
[1] 2
...
[[4]]
[1] 4
> as.numeric("123")
[1] 123
```

## ❑ Data Types – Testing and Coercing Functions

Type	Testing	Coercing
Array	<code>is.array()</code>	<code>as.array()</code>
Character	<code>is.character()</code>	<code>as.character()</code>
Dataframe	<code>is.data.frame()</code>	<code>as.data.frame()</code>
Factor	<code>is.factor()</code>	<code>as.factor()</code>
List	<code>is.list()</code>	<code>as.list()</code>
Logical	<code>is.logical()</code>	<code>as.logical()</code>
Matrix	<code>is.matrix()</code>	<code>as.matrix()</code>
Numeric	<code>is.numeric()</code>	<code>as.numeric()</code>
Vector	<code>is.vector()</code>	<code>as.vector()</code>

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - **Sorting**
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Sorting: sort() vs. order() vs. rank() Function

- Sorting: sort() vs. order() vs. rank() Function

- sort() sorts the vector in an ascending order.

- rank() ranks the numbers in the vector the smallest number receiving the rank 1

- order() returns the indices of the vector in a sorted order

```
> # Build a tibble (a simple data frame)
> set.seed(12345)
> t <- tibble(
+   unit = LETTERS[1:3],
+   a = rnorm(3),
+   b = rnorm(3)
+ )
> class(t)
[1] "tbl_df"      "tbl"        "data.frame"
>
> # Create two new columns: rank_a and rank_b, which, as the names imply,
> # contain the rank (or order) of each value in their corresponding columns.
> t$sort_a <- sort(t$a)
> t$sort_b <- sort(t$b)
> t$order_a <- order(t$a)
> t$order_b <- order(t$b)
> t$rank_a <- rank(t$a)
> t$rank_b <- rank(t$b)
> t
# A tibble: 3 x 9
  unit      a      b sort_a sort_b order_a order_b rank_a rank_b
<chr> <dbl> <dbl> <dbl> <dbl> <int> <int> <dbl> <dbl>
1 A      0.586 -0.453 -0.109 -1.82      3      3      2      2
2 B      0.709  0.606  0.586 -0.453      1      1      3      3
3 C     -0.109 -1.82   0.709  0.606      2      2      1      1
```

## ❑ Data Preparation – Sorting a Vector

- The `sort()` function takes one vector argument, either numeric or character, and returns a vector of sorted values; To sort in decreasing order: `rev(sort(x))`

Example: `> x <- c(1,2,3,2,3,4,8,12,43,-4,-1,NA); sort(x)`  
[1] -4.0 -1.0 1.0 2.0 2.3 3.0 4.0 8.0 12.0 43.0

- The `order()` function sorts a vector, matrix or data frame: `order(x, decreasing = FALSE, na.last = NA, ...)`

- x: vector

- decreasing: decrease or not

- na.last: if TRUE, NAs are put at last position, FALSE at first, if NA, remove them (default)

Example: `> order(x)`

[1] -4.0 -1.0 1.0 2.0 2.3 3.0 4.0 8.0 12.0 43.0

## ❑ Data Preparation – Sorting a Vector

- Example: Sort Vector Continue:

```
> x <- c(1,2.3,2,3,4,8,12,43,-4,-1,NA)
```

```
> order(x)
```

```
[1] -4.0 -1.0 1.0 2.0 2.3 3.0 4.0 8.0 12.0 43.0
```

```
> order(x, decreasing=TRUE)
```

```
[1] 43.0 12.0 8.0 4.0 3.0 2.3 2.0 1.0 -1.0 -4.0
```

```
> order(x, decreasing=TRUE, na.last=TRUE)
```

```
[1] 43.0 12.0 8.0 4.0 3.0 2.3 2.0 1.0 -1.0 -4.0 NA
```

```
> order(x, decreasing=TRUE, na.last=FALSE)
```

```
[1] NA 43.0 12.0 8.0 4.0 3.0 2.3 2.0 1.0 -1.0 -4.0
```

## ❑ Data Preparation – Sorting a Dataframe

- Example: Sort a Dataframe:

```
> library(readr)
> setwd("C:/TeachingUChicago/Spring2020/Data")
> Autodata <- read_csv("AutoCollision.csv")
> str(Autodata)    # Gives the structure of data
> View(Autodata)   # View imported data

# Sort by ClaimSeverity, ascending
> Autodata <- Autodata[order(Autodata$ClaimSeverity),]

# Sort by AgeGroup, descending
> Autodata <- Autodata[order(Autodata$AgeGroup, decreasing=TRUE),]
```

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library



## ❑ Data Preparation – Duplicates

- **Duplicates:**

The function `unique()` will return a dataframe with the duplicate rows or columns removed.

NOTE: `unique()` only work for imported dataframes and doesn't work for dataframes created during an R session

Example:

```
> Dupdata <- read_csv("Dupautocoll.csv")    # 35obs. of 4 variables
> NoDupdata<-unique(Dupdata)                # Dataset with 3 duplicated rows removed: 32obs. of 4 variables
```

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Missing Data

- R denotes data that is not available by *NA*. Quantities that are not a number, such as 0/0, are denoted by **NaN**. In R NaN implies NA (NaN refers to unavailable numeric data and NA refers to any type of unavailable data)
- How a function handles missing data depends on the function. For example, *mean* only ignores *NA*s if the argument *na.rm*=TRUE, whereas which always ignores missing data.

```
> x <- c(4, 7, 2, 0, 1, NA); > mean(x)      [1] NA
```

```
> mean(x, na.rm=TRUE)                    [1] 2.8
```

```
> which(x>4)                             [1] 2
```

- Undefined or null objects are denoted in R by NULL. This is useful when we do not want to add row labels to a matrix.

Example: `> x <- matrix(1:6, ncol=2, dimnames=list(NULL, c("c.1", "c.2"))); x`

	c.1	c.2
[1,]	1	4
[2,]	2	5
[3,]	3	6

## ❑ Data Preparation – Detecting Missing Data

- To test for missing data avoid using identical() and never use ==.; Using identical() relies on unreliable internal computations and “==” will always evaluate to NA or NaN.
- Functions used for detecting missing data,
  - is.na(x) Tests for NA or NaN data in x
  - is.nan(x) Tests for NaN data in x
  - is.null(x) Tests if x is NULL

Example:

```
> x <- c(4, 7, 2, 0, 1, NA) ; x==NA [1] NA NA NA NA NA NA
> is.na(x) [1] FALSE FALSE FALSE FALSE FALSE TRUE
> any(is.na(x)) [1] TRUE
> (y <- x/0) [1] Inf Inf Inf NaN Inf NA
> is.nan(y) [1] FALSE FALSE FALSE TRUE FALSE FALSE
```

## ❑ Data Preparation – Missing Data

- Use `na.omit()` to remove missing data from a dataset
- Use `na.fail()` to signal an error if a dataset contains NA
- Use `complete.cases()` returns a logical vector indicating which rows have no missing data

Example: `> data <- data.frame(x=c(1,2,3), y=c(5, NA, 8))`

```
> na.omit(data) # Remove all rows with missing data
```

```
# Use na.fail to test if a dataset is complete
```

```
NeedCompleteData <- function(data) {
```

```
  na.fail(data) # Return an error message if missing data
```

```
  lm(y, x, data=data)
```

```
}
```

```
NeedCompleteData(data)
```

```
sum(complete.cases(data)) # Get the number of complete cases
```

```
sum(!complete.cases(data)) # Get the number of incomplete cases
```

## Week 3 Class Agenda

- Data Preparation
  - Subsetting
  - Merging
  - Reviewing Data Types
  - Sorting
  - Duplicates
  - Missing values
  - Dplyr library

## ❑ Data Preparation – Dplyr Library

- For analytics, the following functions from dplyr are useful for analytics – Let us see details in the related R scripts
  - 1. filter() (and slice())
  - 2. arrange()
  - 3. select() (and rename())
  - 4. distinct()
  - 5. mutate() (and transmute())
  - 6. summarise()
  - 7. sample\_n() and sample\_frac()



Q & A







**Thank You**





**Contact Information:**

**Your comments and questions are valued and encouraged.**

**Mei Najim**

**E-mail: [mnajim@uchicago.edu](mailto:mnajim@uchicago.edu)**

**LinkedIn: <https://www.linkedin.com/in/meinajim/>**

